Name: Sakshi Anil Mhatre
Roll No: 13
Experiment No.3
Apply Stochastic Gradient Descent algorithm on a feed forward neural network for Iris Flower classification
Date of Performance:
Date of Submission:

Aim: Apply Stochastic Gradient Descent algorithm on a feed forward neural network for Iris Flower classification.

Objective: Ability to perform optimization technique on a feed forward neural network.

Theory:

Gradient Descent is an iterative optimization process that searches for an objective function's optimum value (Minimum/Maximum). It is one of the most used methods for changing a model's parameters in order to reduce a cost function in machine learning projects.

The primary goal of gradient descent is to identify the model parameters that provide the maximum accuracy on both training and test datasets. In gradient descent, the gradient is a vector pointing in the general direction of the function's steepest rise at a particular point. The algorithm might gradually drop towards lower values of the function by moving in the opposite direction of the gradient, until reaching the minimum of the function.

Types of Gradient Descent:

Typically, there are three types of Gradient Descent:

- Batch Gradient Descent
- Stochastic Gradient Descent
- Mini-batch Gradient Descent

Stochastic Gradient Descent (SGD):

Stochastic Gradient Descent (SGD) is a variant of the Gradient Descent algorithm that is used for optimizing machine learning models. It addresses the computational inefficiency of

traditional Gradient Descent methods when dealing with large datasets in machine learning projects.

In SGD, instead of using the entire dataset for each iteration, only a single random training example (or a small batch) is selected to calculate the gradient and update the model parameters. This random selection introduces randomness into the optimization process, hence the term "stochastic" in stochastic Gradient Descent

The advantage of using SGD is its computational efficiency, especially when dealing with large datasets. By using a single example or a small batch, the computational cost per iteration is significantly reduced compared to traditional Gradient Descent methods that require processing the entire dataset.

Stochastic Gradient Descent Algorithm

Initialization: Randomly initialize the parameters of the model.

Set Parameters: Determine the number of iterations and the learning rate (alpha) for updating the parameters.

Stochastic Gradient Descent Loop: Repeat the following steps until the model converges or reaches the maximum number of iterations:

a. Shuffle the training dataset to introduce randomness.

b. Iterate over each training example (or a small batch) in the shuffled order.

c. Compute the gradient of the cost function with respect to the model parameters using the current training.

d. Update the model parameters by taking a step in the direction of the negative gradient, scaled by the learning rate.

e. Evaluate the convergence criteria, such as the difference in the cost function between iterations of the gradient.

Return Optimized Parameters: Once the convergence criteria are met or the maximum number of iterations is reached, return the optimized model parameters.

In SGD, since only one sample from the dataset is chosen at random for each iteration, the path taken by the algorithm to reach the minima is usually noisier than your typical Gradient Descent algorithm. But that doesn't matter all that much because the path taken by the algorithm does not matter, as long as we reach the minimum and with a significantly shorter training time.

Code:

```
#importing the libraries

import numpy as np

import pandas as pd

import tensorflow as tf

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.datasets import load_iris


#Loading and preprocessing the data

iris = load_iris()

X = iris.data

y = iris.target

scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)


#Building the Neural Network Model

model = tf.keras.models.Sequential([

    tf.keras.layers.Input(shape=(4,)),

    tf.keras.layers.Dense(64, activation='relu'),

    tf.keras.layers.Dense(32, activation='relu'),

    tf.keras.layers.Dense(3, activation='softmax')

])


#Compiling the Model

model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics=['accuracy'])


#Training the Model

batch_size = 32

epochs = 100

model.fit(X_train, y_train, batch_size=batch_size, epochs=epochs, validation_split=0.1)


#Evaluating the Model
```

```
test_loss, test_accuracy = model.evaluate(X_test, y_test)

print(f"Test Loss: {test_loss:.4f}, Test Accuracy: {test_accuracy:.4f}")
```

Output:

```
1/1 [==============================] - 0s 118ms/step
Test Accuracy: 0.93
```

Conclusion:

An accuracy of 93% is achieved which indicates that the neural network has learned well. Its structure includes an input layer with four neurons, one or more hidden layers, and an output layer with three neurons.